



# Unit I - Introduction



Network Information System Technologies



# Goals

---

- ▶ Networked systems are Distributed Systems
- ▶ Get an understanding of
  - ▶ What is a distributed system
  - ▶ Why are they relevant
  - ▶ What are the main applications
- ▶ Additionally
  - ▶ Explore some examples
  - ▶ Get a good grasp on the important topic of Cloud Computing



# Index

---

1. Concept of Distributed System
2. Relevance
3. Application Areas
4. Back to the beginning: Cloud Computing
5. Programming Paradigms
6. Conclusions: No computing without a Network



# I. What is a Distributed System

- ▶ Set of autonomous agents
  - ▶ Each agent is a sequential process, proceeding at its own pace.
- ▶ Agents interact. Options:
  - ▶ Message passing → *Better*
  - ▶ Shared memory → *RACE CONDITIONS !!!*
- ▶ Agents have their own independent state
- ▶ There is some collective goal to this cooperation
  - ▶ By which the behaviour of the “system” can be assessed.
- ▶ In practice, a Distributed System is a Networked System.



# Index

---

1. Concept of Distributed System
2. Relevance
3. Application Areas
4. Back to the beginning: Cloud Computing
5. Programming Paradigms
6. Conclusions: No computing without a Network



## 2. Relevance

- ▶ Evolving field since its beginnings
  - ▶ Offshoot of concurrent systems
    - ▶ Heavily studied for their usefulness in the design of time sharing systems
    - ▶ You should be familiar with many aspects of concurrent systems (CSD)
- ▶ Pushed by evolution of computer networks
  - ▶ How to make all those computers do something globally useful?

*Every distributed sys. are Concurrent*

*BUT*

*There are concurrent sys. that are nt distributed*

## 2. Relevance

### ▶ Main “reasons” (as presented back in the 80s)

#### 1. Speed up

*THROUGHPUT INCREASES*  
*Depends on operation →*  
*For READ accesses → throughput increases proportionally to the n° of servers*  
*For UPDATES → throughput cannot be increased linearly with additional cases*

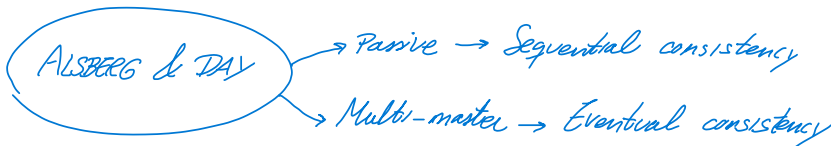
- ▶ Take a complex problem, split it in pieces, have each piece taken care of by a different computer.

#### 2. Fault tolerance. Basic idea:

- ▶ If one computer breaks down, we still have other computers capable of carrying out the tasks of the crashed one.

#### 3. Resource Sharing

- ▶ One computer may have resources (e.g., printers, disks, ...) other computers do not have (and do not need to have)
- ▶ It should be possible to access resources from everywhere





## 2. Relevance

- ▶ All the previous reasons are still valid today because the computing environment IS distributed and interconnected
  - ▶ Myriad of connected “computers”
  - ▶ Myriad of remote services
    - ▶ Accessed as shared resources
    - ▶ Everyone knows and uses the WWW
- ▶ Challenges
  - ▶ Leverage the connectivity to achieve useful results
  - ▶ Create subsystems capable of delivering well-behaved services
    - ▶ How does Google manage to deliver their search engine?
    - ▶ How does Dropbox manage to serve millions of users shared files
    - ▶ ... (your favourite service goes here)





# Index

---

1. Concept of Distributed System
2. Relevance
3. Application Areas
4. Back to the beginning: Cloud Computing
5. Programming Paradigms
6. Conclusions: No computing without a Network



## 3. Application Areas

- Some important <sup>DISTRIBUTED</sup> application areas are:
1. *World Wide Web*
  2. *Sensor networks*
  3. *Internet of Things*
  4. *Cooperative computing*
  5. *Highly available Clusters*
- They are described on the sequel...



## 3. 1. Application Areas: The WWW

---

- ▶ Based on the Client-Server Model
- ▶ Server attends requests for documents
  - ▶ Requests may involve reading or writing of a document
- ▶ The Clients are Browsers, sending/receiving documents
  - ▶ Browsers parse documents searching for metadata
  - ▶ Links are particular metadata pointing to other documents
    - ▶ Documents may be in another server
- ▶ Simple and powerful paradigm
  - ▶ Initially conceived for document sharing
  - ▶ Extended to allow document requests to stand for general service requests
    - ▶ Returned “documents” encode the result of the actual request



## 3. 2. Application Areas: Sensor Networks

---

- ▶ Driven by declining costs of hardware
- ▶ Special purpose mini-computers
  - ▶ Motes
- ▶ Embedded in common devices
  - ▶ Dishwashers, etc
- ▶ Contain physical world sensors
  - ▶ Humidity, temperature, power consumption, ...
- ▶ Wide range of potential applications
  - ▶ Surveillance
  - ▶ Biological and chemical disaster detection
  - ▶ Power monitoring
  - ▶ ...



## 3. 3. Application Areas: The “Internet of Things”

---

- ▶ Motivation: leverage ubiquitous connectivity of all devices
  - ▶ Generalization of sensor networks
    - ▶ All devices can, and will interact among them
    - ▶ Devices can also alter their physical environment
  - ▶ New scenarios open up
    - ▶ Smart cities
    - ▶ Building automation
    - ▶ Healthcare
    - ▶ ...

### 3.3. Application Areas: The “Internet of Things”





## 3.4. Application Areas: Cooperative Computing

---

- ▶ Most computational power is underused
  - ▶ The desktops spend many hours doing nothing
- ▶ Many engineering and scientific problems can be split into pieces (tasks)
  - ▶ Each task can be resolved in a small amount of time
  - ▶ Results from each piece can be composed to complete the resolution of the whole problem
- ▶ Servers can be set up with an instance of such a problem
  - ▶ The server creates a pool of smaller tasks
- ▶ Computers across the internet can subscribe to receive tasks to solve
  - ▶ They install a special client software: the task runtime environment
  - ▶ The client registers with the server
- ▶ The server spreads tasks among the registered clients, and collects their results



## 3. 5.Application Areas: Highly Available Clusters

---

- ▶ So far we have seen application areas addressing resource sharing and cooperation.
- ▶ Fact:
  - ▶ Devices fail. Computers are devices. They fail at some point with a 100% probability.
- ▶ Fact:
  - ▶ Not all devices fail at the same time, always.
  - ▶ Q: when can it happen?
- ▶ Some environments need a high degree of availability
  - ▶ Banking
  - ▶ Finances
  - ▶ ...
- ▶ Leverage having more than one device to stand failures





## 3.5. Application Areas: Highly Available Clusters

---

### ▶ HA Cluster:

- ▶ Set of computers, with server programs on which clients depend constantly
- ▶ Typically holding sensitive data
- ▶ Designed with specific protocols to stand failures of one or more of them
- ▶ Two main concerns:
  - ▶ Preserve data integrity
  - ▶ Preserve server operation availability



## 3.6.Application Areas: Cloud Computing

---

- ▶ Main current trend to build and provide services
- ▶ Fact:
  - ▶ Computer power in the traditional usage is underused
    - ▶ We already discussed this earlier
  - ▶ Setting up enterprise-class computer centres, with their applications is expensive:
    - ▶ Purchasing the software and hardware
    - ▶ Paying up the engineers which manage hardware and software
    - ▶ Paying for energy consumption
      - Onerous, when resources are underutilized

... Let us dive into CC...



# Index

---

1. Concept of Distributed System
2. Relevance
3. Application Areas
4. Back to the beginning: Cloud Computing
5. Programming Paradigms
6. Conclusions: No computing without a Network



## 4. CC: Cloud Computing

---

- ▶ Items to be discussed:
  1. Programs and services
  2. Roles in the service life cycle
  3. Evolution of software services
    - a) Mainframes
    - b) Personal computers
    - c) Enterprise computer centres
    - d) SaaS
    - e) IaaS
    - f) SaaS on IaaS
    - g) PaaS
  4. Summary



## 4.1. Software and Services

---

- ▶ **Generic Goal of CC:**
  - ▶ Make creation and exploitation of services based on software simpler and more efficient
- ▶ **An obvious fact:**
  - ▶ Software has always been made to get some sort of service
  - ▶ Helped by hardware, of course
- ▶ **Computer industry evolution helped to obscure this fact:**
  - ▶ Personal Computer industry has imposed a specific mode of interaction of users with their computers.



## 4.2. Roles in the Service Life Cycle

---

- ▶ Consider the following four roles:
  - ▶ **Developer**
    - ▶ Gets the software components built
  - ▶ **Service provider**
    - ▶ Decides the characteristics of a service, the components that make it up, and how it should be configured and managed
  - ▶ **System's administrator**
    - ▶ Makes sure every piece of software/hardware is in place and properly configured
  - ▶ **Service user**
    - ▶ Accesses the service



## 4.3. Service Evolution: a) Mainframes

---

- ▶ *System administration* taken care of by specialists
- ▶ Very few tensions of contention
  - ▶ Small enough user base
- ▶ Efficient use of hardware
  - ▶ Shared by a large population of users
  - ▶ Low up-front cost for a user
    - ▶ Purchase cost born by Mainframe owner
- ▶ Mixed role for Users
  - ▶ Many were also the developers
  - ▶ Many were also their own service providers
    - ▶ With the software they developed
    - ▶ With third party software
- ▶ Users were involved in too many of the management details for the services they finally used



## 4.3. Service Evolution:

### b) Personal Computers/Workstations

---

- ▶ Trends with increased computer power produced the personal computer
  - ▶ Users no longer needed to have access to a mainframe in a computer centre
- ▶ No contention
  - ▶ Precisely one of the selling points of the model
- ▶ Wasteful use of resources: computer infra-utilized
- ▶ Up-front investment cost to purchase
- ▶ Rationalization of the role of developer
  - ▶ Specialized organizations build the software
- ▶ But still, mixed role of user
  - ▶ Works as service provider
    - ▶ Selects the mix of software she needs to get some job done
  - ▶ Works as system administrator of her PC
- ▶ Too much complexity for the majority of users





## 4.3. Service Evolution:

### c) Enterprise Computer Centers

---

- ▶ Including HA Cluster set-ups
- ▶ Similar characteristics to the PC situation
  - ▶ The user being the enterprise
  - ▶ High Personnel cost to take up the roles of system administrator and service provider
  - ▶ On occasion, developer role, for in-house software
- ▶ Variation based on hosting the software on external Data Centre
  - ▶ Avoid up front hardware costs
  - ▶ Avoid costs of hardware maintenance and management
  - ▶ Avoid constant power usage costs
  - ▶ Better able to manage computing costs



## 4.3. Service Evolution:

### d) Software as a Service (SaaS)

---

- ▶ Service is accessed through the network.
  - ▶ Typically using a web browser
- ▶ Clear separation of the role of user
  - ▶ Service is defined by a third party: the service provider
- ▶ Not so clear separation of the other roles
  - ▶ Software is initially developed mostly by the provider
  - ▶ All management aspects are taken care of by the provider
    - ▶ Including management of hardware in data centres
    - ▶ Including management of the installed software on the hardware
- ▶ Initially some inefficiencies
  - ▶ Lack of flexibility in hardware allocation
    - ▶ Leads to provider committing to a certain resource usage.
    - ▶ Diminishes resource sharing
  - ▶ Limited contention: resource reservation for expected demand



## 4.3. Service Evolution: d) SaaS

---

- ▶ Driven by
  - ▶ Improvement of networking technology
    - ▶ Higher bandwidth
    - ▶ Lower latency
  - ▶ Built capacity of Data Centres
    - ▶ Made it attractive to host services for external users
  - ▶ Technology improvements in the Browser
    - ▶ Epitomized by the buzzword “Web 2.0”
    - ▶ Browser capable of executing locally complex user interactions, delivered via Browser-executing software
      - Allowed attractive user interfaces
      - Improved scalability:
        - Workload reduction on servers.



## 4.3. Service Evolution:

### e) Infrastructure as a Service (IaaS)

---

- ▶ Provides ability to easily allocate/de-allocate computer and network resources on demand
  - ▶ Requests via API to a service (the IaaS service)
  - ▶ Ability to load custom OS images on those computers
  - ▶ Ability to request concrete computer and network capacities
- ▶ Driven by Hardware Virtualization technology
  - ▶ Easy and fast allocation/deallocation of hardware resources (virtual)
  - ▶ Easy sizing of hardware resources (virtual)
  - ▶ Easy set-up of a computer image within a virtual machine



## 4.3. Service Evolution: f) SaaS on IaaS

- ▶ IaaS introduces pay-as-you-go model
  - ▶ A core characteristic of Cloud Computing
- ▶ Makes it feasible to create SaaS which adapt to their user's load
  - ▶ The more user load, the more hardware resources it requests
    - ▶ *Elasticity*: another Cloud characteristic
  - ▶ Moves the pay-as-you-go model to the SaaS
    - ▶ Users of SaaS also pay per use of the service
- ▶ Very efficient usage of resources for SaaS providers
  - ▶ Most costs are variable
  - ▶ No up-front costs to reserve capacity (purchase or commitments)
  - ▶ Savings passed to SaaS users: competitive market of services.



## 4.3. Service Evolution: f) SaaS on IaaS

---

- ▶ IaaS providers take the risk of up-front investment
  - ▶ Promise of a large population of SaaS providers
    - ▶ Themselves with large populations of SaaS users
    - ▶ Large demand of virtualized resources
- ▶ SaaS provider still has mixed roles
  - ▶ Software Service provider (its natural role)
  - ▶ Must manage allocation of hardware resources
  - ▶ Must manage OS images, their upgrades, and their base software
  - ▶ Must build their own service management strategy
    - ▶ Monitoring mechanisms
    - ▶ Upgrade mechanisms




## 4.3. Service Evolution:

### g) Platform as a Service (PaaS)

---

- ▶ Ideally, promises to take away extraneous tasks from SaaS providers
  - ▶ Still in infancy, though
- ▶ Purports to be the equivalent of an OS
  - ▶ Specifies a Service Model on which to base specification of SaaS, and development of their software components.
  - ▶ Includes the following aspects
    - ▶ Configuration and lifetime management model (including dependency expressions)
      - Composition, Configuration, Deployment and Upgrade mechanisms
    - ▶ Performance model
      - Automatic monitoring of relevant parameters
      - Expression of elasticity points
      - Automatic reconfiguration under varying load

- ▶ Cloud Computing: Ideal layering







## 4.4. Summary

---

- ▶ Cloud Computing is all about efficiency and easiness:
  - ▶ Efficient sharing of resources
    - ▶ Consume only what one needs
    - ▶ Pay only for what is used
  - ▶ Easy adaptation to a varying population of users
  - ▶ Easy ways to develop and provide a service
- ▶ Three layers of cloud services are identified:
  - ▶ Software as a Service
    - ▶ The actual goal is to provide these to a large user population
  - ▶ Platform as a Service
    - ▶ Extremely desirable to automate management of resources for SaaS and ease creation and deployment of services.
  - ▶ Infrastructure as a Service
    - ▶ Enables SaaS elasticity
- ▶ From the user's point of view CC feels like going back to the Mainframe era



# Index

---

1. Concept of Distributed System
2. Relevance
3. Application Areas
4. Back to the beginning: Cloud Computing
5. Programming Paradigms
6. Conclusions: No computing without a Network



## 5. Programming Paradigms

---

- ▶ A prevalent way to organize a distributed system is to make each process a “server”
  - ▶ Receives requests, processes them, sends back answers
- ▶ Servers, themselves, request services from other servers
  - ▶ They may need such services to satisfy a request they receive
- ▶ To be able to scale, servers must not BLOCK while serving request
  - ▶ Should be able to accept other requests
- ▶ Two programming paradigms exist:
  - ▶ Multi-threaded programming
    - ▶ State-sharing concurrent servers
  - ▶ Asynchronous programming
    - ▶ Asynchronous single-threaded servers



## 5.1. Concurrent, state-sharing servers

---

- ▶ Multi-threaded concurrent programs
  - ▶ Each request is handled in its own thread
  - ▶ All threads share a global state
  - ▶ Concurrency-control mechanisms are used to implement atomicity
- ▶ Advantages
  - ▶ Threads may block waiting for requests from the server to complete, without blocking the server
- ▶ Disadvantages
  - ▶ Multi-threaded programming has its own overheads
    - ▶ Need to support concurrency control constructs
  - ▶ It turns out that shared memory concurrent programming is
    - ▶ Hard to do well
    - ▶ Hard to reason about how it works
- ▶ Prevalent Environments:
  - ▶ Java
  - ▶ .NET



## 5.2. Async Servers

- ▶ Async programming *aka* event-driven programming.
  - ▶ Closely matches the guard-action program model
  - ▶ Async programs have many activities, but...
    - ▶ State is never concurrently shared among the executing activities
- ▶ “Events” are the “guards”.
- ▶ Actions are established as callbacks of events
  - ▶ Need to dynamically built actions/guards to facilitate programming
  - ▶ When building actions, programming language mechanisms are used to easily establish the state to be affected by the action
    - ▶ Reduce complexity of “preparing” state to link internal actions
- ▶ Actions ready for execution placed on a “turn queue”
  - ▶ Scheduled actions executed in FIFO order of the queue



## 5.2. Async Servers

---

### ▶ Advantages

- ▶ Shared state handling complexity greatly diminishes
  - ▶ Still, careful on handling of the turn queue to avoid surprises
- ▶ Less overhead, as no multi-threading environment is supported
  - ▶ Better ability to scale
- ▶ Close match to how a distributed system actually work: event-driven
  - ▶ Easier to reason about what is going on

### ▶ Disadvantages

- ▶ Proper state-handling is necessary when building actions
- ▶ Needs all environment to be async, not just IPC
  - ▶ OS services need to be async too, to avoid blocking
- ▶ Prevalent environments with built-in support in language
  - ▶ Nodejs
  - ▶ Async .NET



# Index

---

1. Concept of Distributed System
2. Relevance
3. Application Areas
4. Back to the beginning: Cloud Computing
5. Programming Paradigms
6. Conclusions: No computing without a Network



## 6. Conclusions

---

- ▶ Networked Systems are Distributed Systems
  - ▶ Most computing nowadays is networked
    - ▶ Thus, distributed
  - ▶ Proper design and implementation requires mastering aspects of concurrent programming and properties of the architecture
- ▶ Rich set of application areas already exploited
- ▶ Important trend of Cloud Computing as a logical endpoint in the evolution of computing
  - ▶ Driven by efficiency in resource usage
  - ▶ Pay-as-you-go model of access
  - ▶ Elasticity and scalability
- ▶ Two programming paradigms for distributed service development:
  - ▶ Concurrent (i.e., multi-threaded) servers
    - ▶ Should deal with race conditions. Threads may be blocked.
  - ▶ Asynchronous servers
    - ▶ Event-driven. Easily scalable.